

Using Google Consumer Surveys

Lars Vilhuber

This document reports a simple experiment using Google Consumer Surveys.

```
# define basic parameters  
price_quoted <- 3  
price_minimum <- 0.1  
budget <- 500  
coupon <- 75  
result_file <- c("Employment and search.xls")
```

Survey instrument

We based our survey instrument on the primary labor force question on the CPS, somewhat modified:


Please complete the following survey to access this premium content.

Question 1 of 2 or fewer:

Last week, did you do ANY work for pay?

- ☐ No, I am retired
- ☐ No, I am disabled
- ☐ No, I was unable to work
- ☐ Yes
- ☐ No

OR

 Show me a different question

 Skip survey

Google

[INFO](#) [PRIVACY](#)

We followed up for all non-employed (all those not responding with “Yes”“) with a question on search intensity that was not based on the CPS:

Question 2 of 2:

If you are not working, how many jobs did you apply for (by email, website, letter, in person, etc.) last week?

☐ 0 (I am not looking for work)


☐ 1-2


☐ 3-5

☐ 6-10

☐ more than 10

OR

 Show me a different question

 Skip survey

Google

[INFO](#) [PRIVACY](#)

Results

Google first ran a test (which resulted in being 4 responses) to assess how much to charge us. Google claims that prices are between \$0.1 and \$3. In our case, the quoted price was \$3 per response, and our budget was \$500 (although we were offered a one-time coupon valued at \$75. We thus requested 191 cases. Once accepted, Google posted the survey, targetting the “general population in the United States on the Google Consumer Survey publisher network”.

```
# get overview data
xls_overview <- read.xls(result_file,sheet=1)
xls_topline <- read.xls(result_file,sheet=2)
xls_data <- read.xls(result_file,sheet=3)
```

To our initial surprise, we had received 461 cases. As it turns out, since our labor force question served as a “screener”, we were only charged for complete questions, i.e., responses that had completed both Q1 and Q2 of which we had obtained 192. Thus, the final response was

Last week, did you do ANY work for pay?

If you are not working, how many jobs did you apply for (by email, website, letter, in person, etc.) last week?

From this, we can compute a few statistics of interest.

```
results <- data.frame(c(0,0,0),row.names = c("Employment-Population Rate","Unemployment rate","OLF rate")
results[,1:2] <- 0
names(results) <- c("Google","BLS")
# define OLF, pop, unemployed, employed
pop <- xls_overview[1,c("Response.count")]
nonemp <- xls_overview[2,c("Response.count")]
employed <- pop - nonemp
olfind <- xls_data$Question..2.Answer=="0 (I am not looking for work)"
olf <- nrow(xls_data[olfind,])
google.dates <- c(
  min(as.POSIXct(strptime(xls_data[, "Time..UTC."], "%Y-%m-%d %H:%M:%S"), tz="UTC")),
  max(as.POSIXct(strptime(xls_data[, "Time..UTC."], "%Y-%m-%d %H:%M:%S"), tz="UTC"))
)
# now compute stats
results["Employment-Population Rate",1] <-
  round(employed*100/pop,2)
results["Unemployment rate",1] <-
  round((nonemp-olf)*100/employed,2)
results["OLF rate",1] <-
  round(olf*100/pop,2)
```

which we measured between 2016-01-08 11:10:46 and 2016-01-08 20:20:05.

For comparison purposes, we obtained the BLS unemployment rate and employment-population ratio from FRED:

```
# define the series of interest
fred_unemployment <- c("UNRATENSA")
fred_unemployment20 <- c("LNU04000024")
fred_emppop_all <- c("LNU02300000")
fred_emppop_men <- c("LNU02300001")
fred_emppop_women <- c("LNU02300002")
fred_olf_count <- c("LNU05000000")
fred_pop_count <- c("CNP160V")

#bls.urate <- get(getSymbols(fred_unemployment,src='FRED'))
bls.urate20 <- get(getSymbols(fred_unemployment20,src='FRED'))
bls.emppop_all <- get(getSymbols(fred_emppop_all,src='FRED'))
#bls.emppop_men <- get(getSymbols(fred_emppop_men,src='FRED'))
#bls.emppop_women <- get(getSymbols(fred_emppop_women,src='FRED'))
#
# Compute the OLF rate
cnt.olf <- get(getSymbols(fred_olf_count,src='FRED'))
cnt.pop <- get(getSymbols(fred_pop_count,src='FRED'))
bls.olf_rate <- cnt.olf*100/cnt.pop
#
# attach the closest value - this might change over time, but that's fine.
```

```

# found at http://stackoverflow.com/questions/8186960/finding-the-most-recent-observation-earlier-than-
# For emppop ratio, use all - we will break it out by gender later
results["Employment-Population Rate",2] <-
  round(bls.emppop_all[findInterval(google.dates[1],index(bls.emppop_all)),]
,2)
results["Unemployment rate",2] <-
  round(bls.urate20[findInterval(google.dates[1],index(bls.urate20)),],2)
results["OLF rate",2] <-
  round(bls.olf_rate[findInterval(google.dates[1],index(bls.olf_rate)),],2)

```

The final (unweighted) result is then as follows (note: the BLS numbers won't add up quite to 100, since they stem from different series):

	Google	BLS
Employment-Population Rate	58.35	59.40
Unemployment rate	16.36	4.50
OLF rate	32.10	37.59

Appendix

```

# libraries needed
library(quantmod) # read FRED
library(gdata) # read Excel
options("getSymbols.warning4.0"=FALSE)
Sys.info()

```

```

##                sysname
##                "Linux"
##                release
##                "3.16.7-29-desktop"
##                version
## "#1 SMP PREEMPT Fri Oct 23 00:46:04 UTC 2015 (6be6a97)"
##                nodename
##                "zotique2"
##                machine
##                "x86_64"
##                login
##                "vilhuber"
##                user
##                "vilhuber"
##                effective_user
##                "vilhuber"

```